

Machine Learning Course - Project 1

Victor Faramond, Mathieu Schopfer, Dario Anongba Varela
Department of Computer Science, EPFL Lausanne, Switzerland

I. INTRODUCTION

II. MODELS AND METHODS

A. Data

Two sets of data were provided in the form of csv files, a *training* and a *test* set. The input of both sets consists in $D = 30$ feature columns \mathbf{x}_n . In the training set, an additional output column y_n is provided. The output is binary (-1 or 1), depending whether the recorded signal corresponds to an actual event (1) or to background noise (-1).

B. Pre-processing

1) *Missing data*: After an exploratory data analysis phase, we found that only a quarter of the dataset is complete. 11 variables out of 30 have missing records which are designated by -999. However, we found a clear structure about these missing values. Most of them are related to the *jet* properties. If an event contains no jet at all, all the jet related features are missing. If an event contains only one jet, the leading jet features are present but the subleading jet features are missing. Therefore, in the ridge regression case (see below), we decided to split the dataset into three different subsets corresponding to the jet values $\{0\}$, $\{1\}$ and $\{2, 3\}$, drop the columns with missing values and train a specific classifier on each subset.

2) *Missing values imputing*: Missing values are first imputed. For a given input column $x_{\bullet,i}$ ($i \in \{1, \dots, D\}$) missing values are replaced in both the training and the test sets by the value appearing most frequently in column $x_{\bullet,i}$ of the training set.

3) *Feature engineering*: For all the positive value columns, we added a corresponding column with the inverse log values of the column,

$$x_{\bullet,i} \rightarrow \frac{1}{1 + \log x_{\bullet,i}}.$$

Thus, we can reduce the range of the values.

4) *Standardization*: We standardized both the training and test datasets using the mean and the variance of the training set.

III. MACHINE LEARNING MODELS

A. Models

The following models were tested:

- Gradient and stochastic gradient descent, $\gamma = 0.01$
- Least squares

- Ridge regression, with polynom basis, $\lambda = 0.6$
- Logistic and regularized logistic regression, $\gamma = 0.6, \lambda = 0.04$

1) *Ridge regression*: Before fitting the ridge regression model, we added polynomial features to our training and test sets:

- The square root $\sqrt{x_n}$ of the feature.
- \mathbf{x}_n^k , for $k \in \{1, \dots, d\}$. The parameter d is different for each subset of data (regarding the number of jets) and has been defined using a grid search with values of $d \in \{1, \dots, 11\}$. See figure 1.
- The products $x_{\bullet,i} \cdot x_{\bullet,j}$, $\forall i, j \in \{1, \dots, D\}$.

The parameter λ was also defined using a grid search with values of $\lambda \in 1e - 9, \dots, 1e + 1$. See figure 2.

B. Cross validation

To train and test the various models presented hereabove, a 10-fold cross validation method was used. This consists in splitting the training data set into ten subsets of equivalent size. Then 9 subsets are used to train the models, the last one to test it. This process is repeated for all 10 combinations of training and testing subsets. This allows us to make sure a model is stable and that we did not overfit on the training set.

C. Assessment of models accuracy

At each step of the 10-fold cross validation, the positives rate was calculated on the test subset. The positive rates is simply given as the number of model outputs y_n agreeing with the (known) expected output divided by the total number of events in the subset.

Then the mean and variance of positives rates was calculated and retained as the model accuracy score.

D. Choice of model for submission

The model that was chosen for submission was that achieving the highest positives rate score.

IV. RESULTS

Table I summarizes the results for all models. Regarding these results, the ridge regression model was selected for submission.

Figure 1. Mean accuracy of the ridge regression model on data subset verifying jet=?, for various values of maximal polynomial degree.

Figure 2. Mean accuracy of the ridge regression model on data subset verifying jet=0, for various values of λ .

Model	Mean accuracy	Variance
Gradient descent	0.763168	0.000006
Stochastic gradient descent	0.686248	0.000562
Least squares	0.776600	0.000007
Rigde regression	0.833556	0.000002
Logistic regression	0.679196	0.000004
Regularized logistic regression	0.679196	0.000004

Table I
MEAN ACCURACY AND STANDARD DEVIATION FOR ALL TESTED MODELS.